

Tutorial on Generative Adversarial Networks - From basics to current state-of-the-art, and towards key applications in medicine

Răzvan V. Marinescu

Medical Vision Group, Massachusetts Institute of Technology



Slides available online: <https://people.csail.mit.edu/razvan>

- ▶ GAN basics
- ▶ State-of-the-art
- ▶ How to use them?
- ▶ Potential applications in medicine

- Generative: can generate new data instances

$$p(X, Y)$$

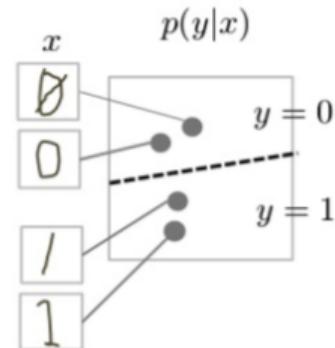
- Discriminative: discriminates between different kinds of data instances

$$p(Y|X)$$

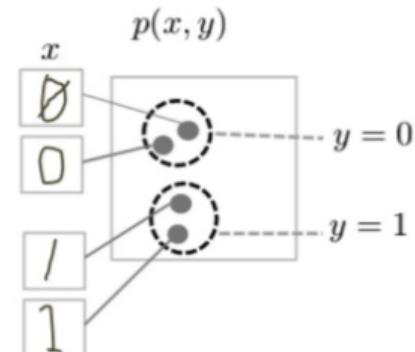
X = image

Y = label/score

- Discriminative Model

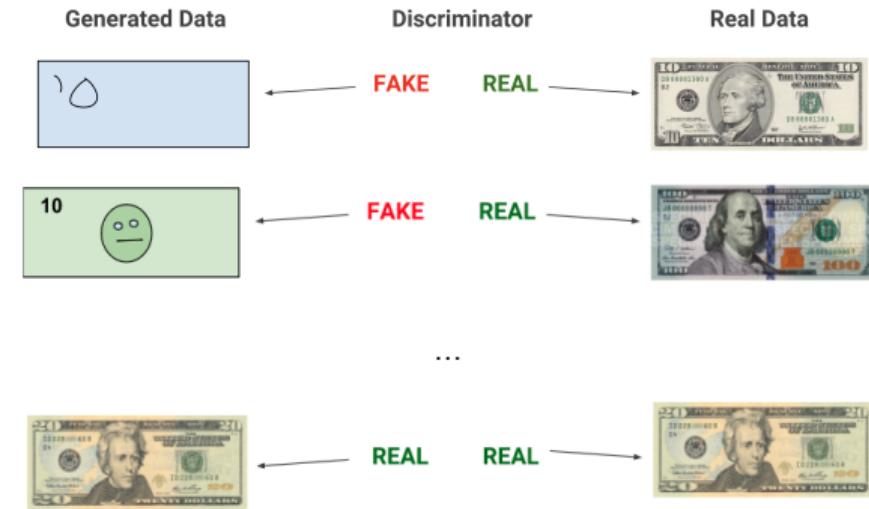


- Generative Model



Introduction to Generative Adversarial Networks

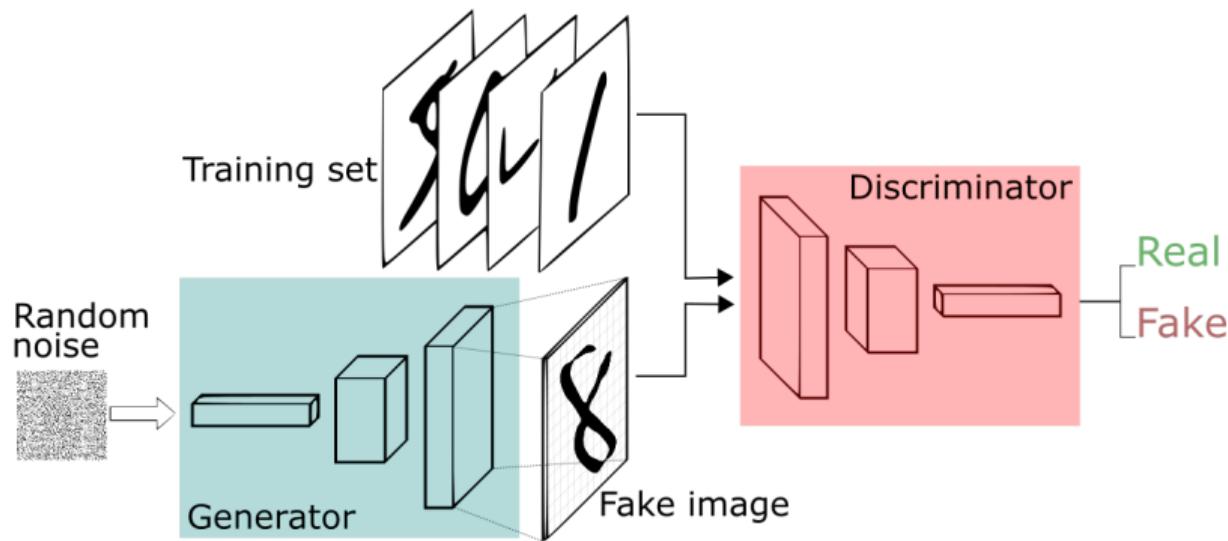
- ▶ Introduced by Goodfellow et al, 2014
- ▶ Two adversaries (generator + discriminator) compete with each other
- ▶ Over time, the generator gets better at generating images



GAN architecture

- Generator attempts to generate good images to fool the discriminator
- Discriminator attempts to tell apart the fake images from the real ones
- Loss function:

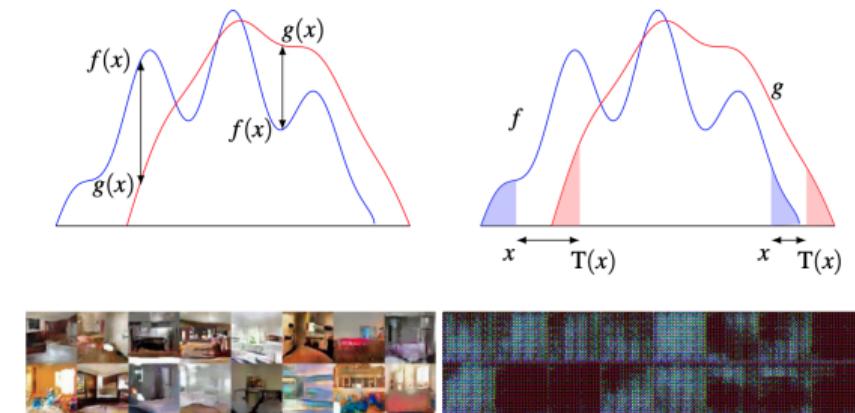
$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$



Towards the state-of-the-art in Generative Adversarial Networks

Wasserstein GANs

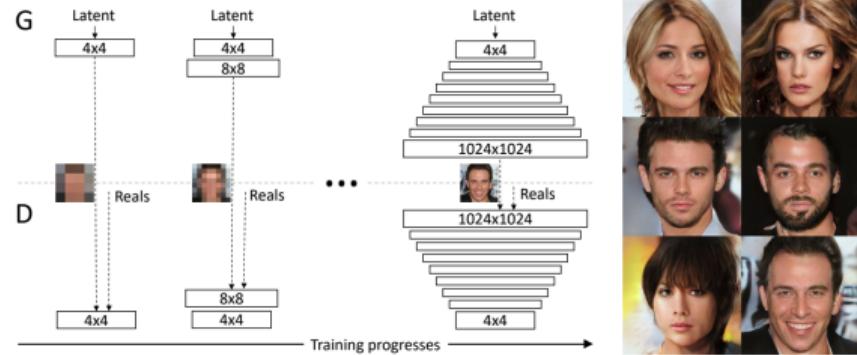
- ▶ Original GANs were very hard to train, often collapsing during training
- ▶ Problem: They optimise the Jensen-Shannon divergence, a “vertical” distance → no good gradients when distributions far away
- ▶ A “horizontal” distance (e.g. Wasserstein) ensures gradients are non-zero when distributions don’t have support (i.e. far away)
- ▶ GAN trained with the new Wasserstein metric collapses less often, and generates better images



(Arjovsky, ICML, 2017)

Progressive Growing of GANs

- ▶ GAN training unstable if one starts directly in high-resolution
- ▶ Key idea: start from low-resolution (4×4) and build up to highest-resolution (1024×1024)
- ▶ Each new layer is faded-in slowly



Mao et al. (2016b) (128 × 128)

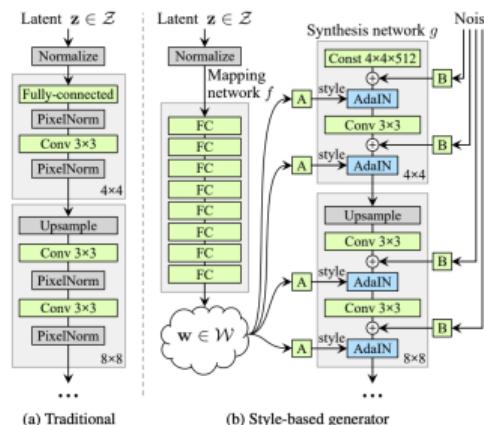
Gulrajani et al. (2017) (128 × 128)

Our (256 × 256)

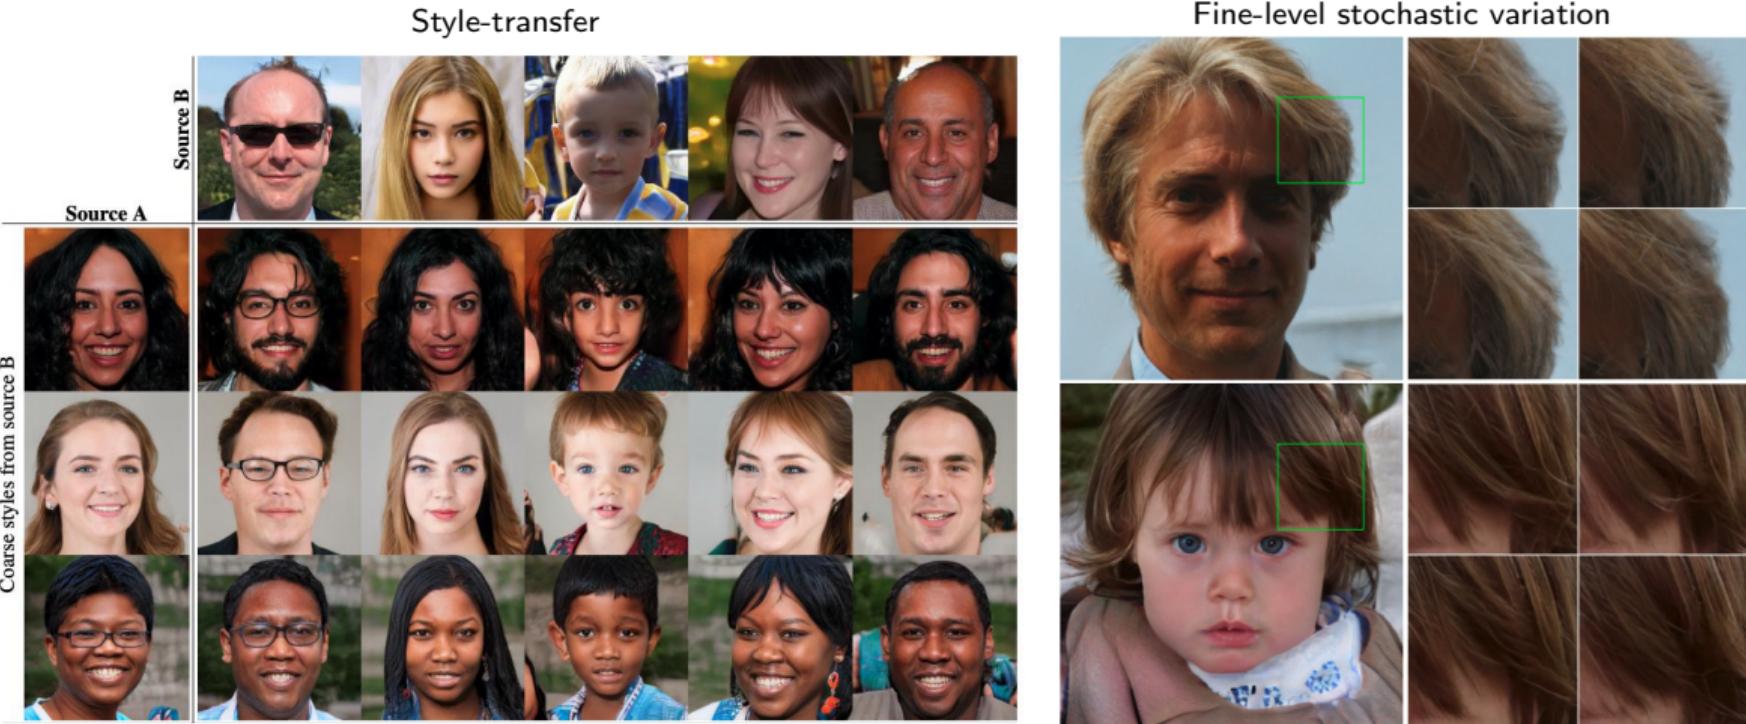
(Karras et al, 2018)

StyleGAN1 (Karras et al, 2019)

- ▶ GAN that borrows ideas from style-transfer literature
- ▶ Uses a mapping network to generate “style vectors” at every level in the generator
- ▶ Each style vector is intensity-normalized (AdaIN operation)
- ▶ Generated images have unprecedented realism and diversity



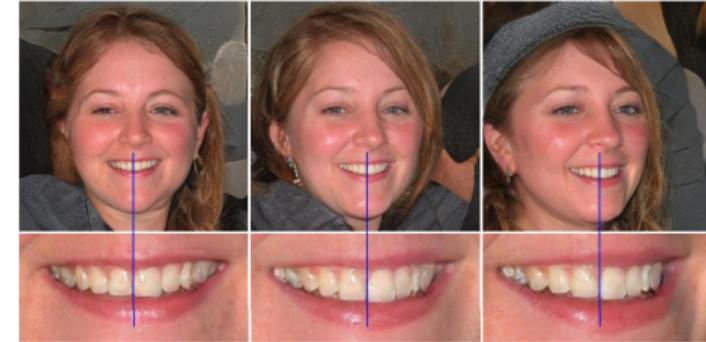
The style-based architecture allows style-transfer along with fine-level stochastic variation



Blob-artefacts caused by AdaIN normalisation



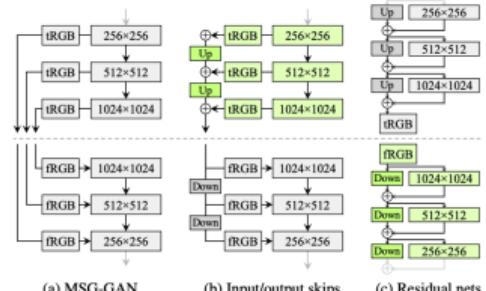
“phase” artefacts due to progressive growing



Solution: bake normalisation straight into convolution weights:

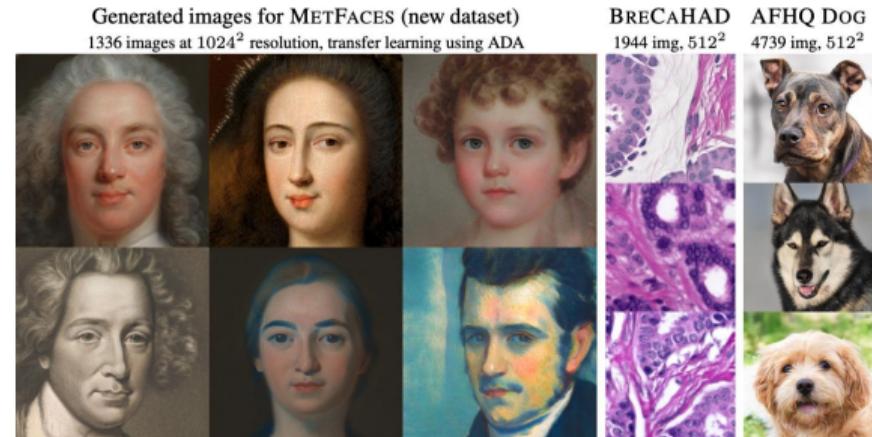
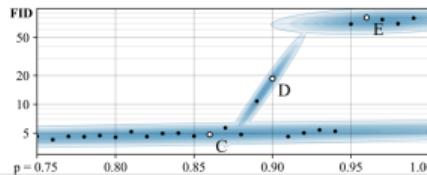
$$w''_{ijk} = w'_{ijk} / \sqrt{\sum_{i,k} {w'}_{ijk}^2 + \epsilon}$$

Solution: communicate across resolution levels through skip connections



Training GANs with limited data (Karras et al, 2020)

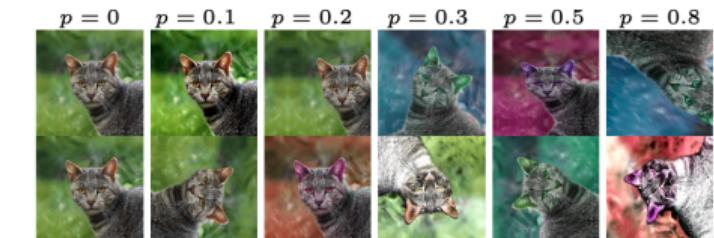
- ▶ Previous StyleGAN2 model needed large number of images for training ($\approx 70,000$)
- ▶ Aim: Enable training on limited datasets (1000 images) through data-augmentation
- ▶ Problem: augmentations leak into the generated images
- ▶ This is mitigated by ensuring augmentation probability p is lower than a threshold (≈ 0.9).



BRECAHAD 1944 img, 512^2 AFHQ DOG 4739 img, 512^2



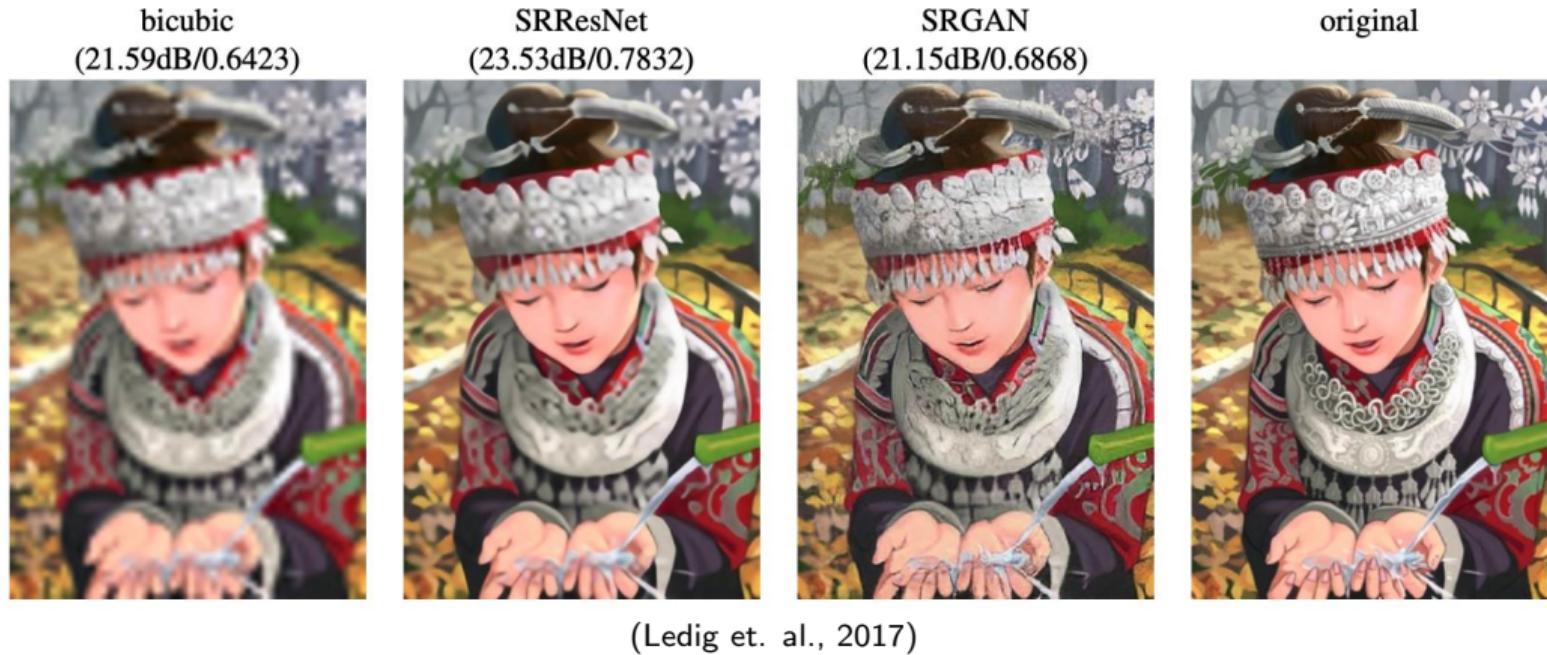
BRECAHAD 1944 img, 512^2 AFHQ DOG 4739 img, 512^2



(c) Effect of augmentation probability p

Applications of GANs and other generative models

Application 1: Image super-resolution



- ▶ G generates high-res image from low-res input
- ▶ D discriminates whether high-res image is fake or real.

Application 2: In-painting



(a) Input context

(b) Human artist



(c) Context Encoder
(L_2 loss)

(d) Context Encoder
($L_2 + \text{Adversarial loss}$)

(Pathak et al, 2016)

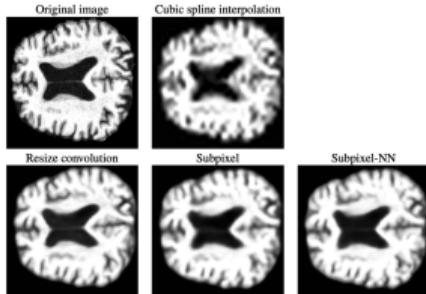
Application 3: Image-to-image translation



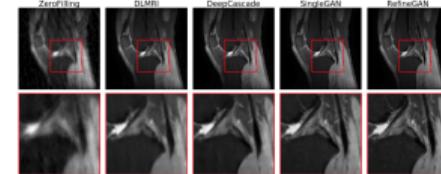
CycleGAN (Jun-Yan Zhu, CVPR, 2016)

Applications in medical imaging

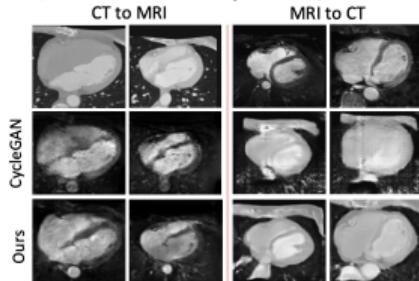
MRI super-resolution (Sanchez et al., 2018)



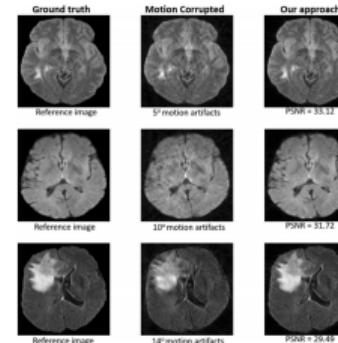
MR Reconstruction from undersampled K-space (Quan et al, 2017, Yang et al, 2017)



Modality translation (Zhang et al, 2019)



MRI motion correction (Usman et al, 2020)



Any image reconstruction task!

Applications of generative models to medicine: prediction of disease progression

- ▶ Prediction and visualisation of future of disease progression
- ▶ Can assist doctors in assigning treatments

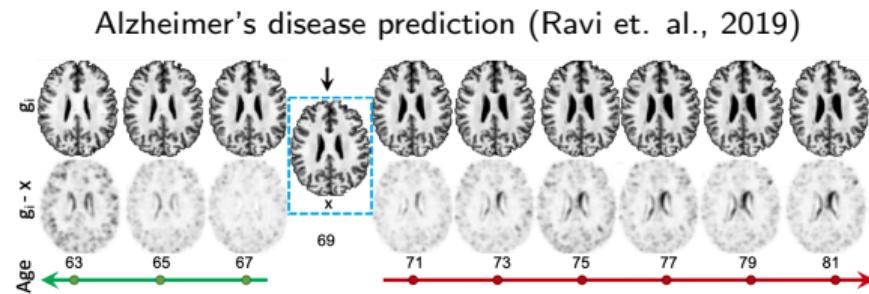
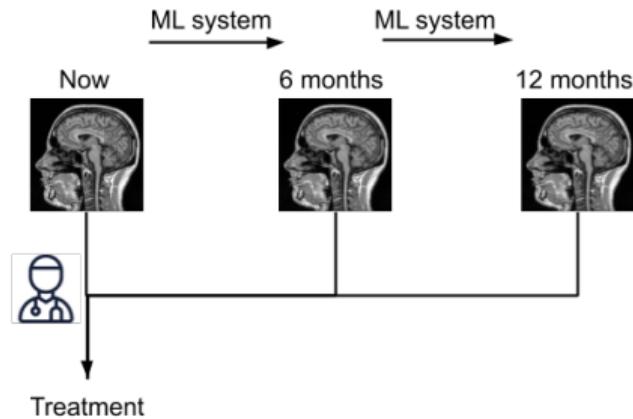
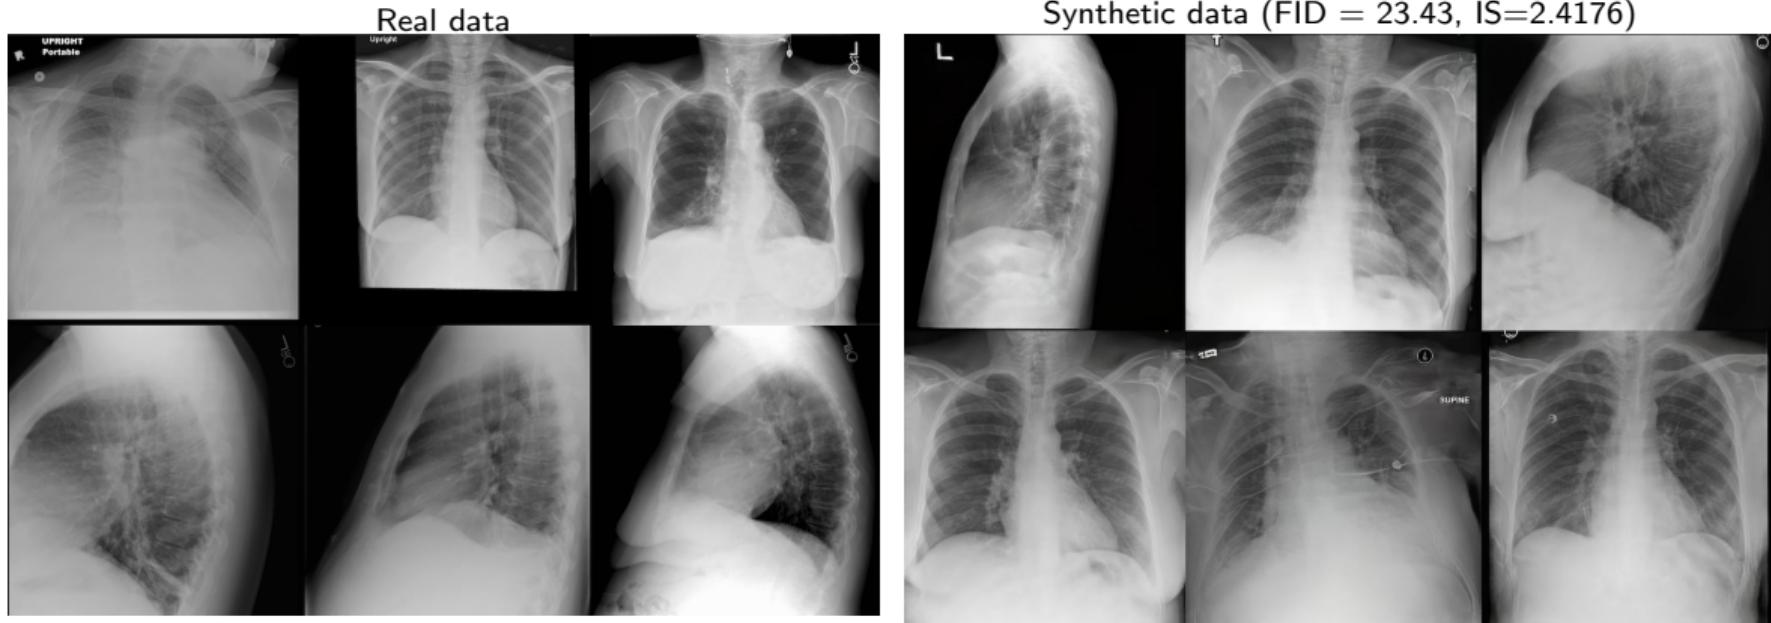


Fig. 4. Neurodegeneration simulation of a 69-year old ADNI participant.

Preliminary results of StyleGAN2 on three medical datasets

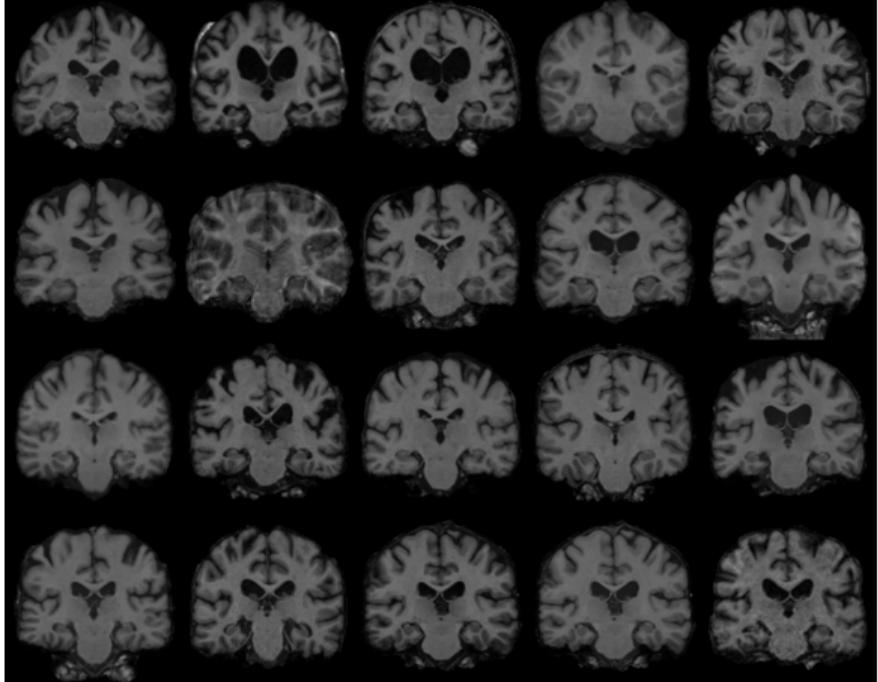
Preliminary results of StyleGAN2 on Chest X-rays



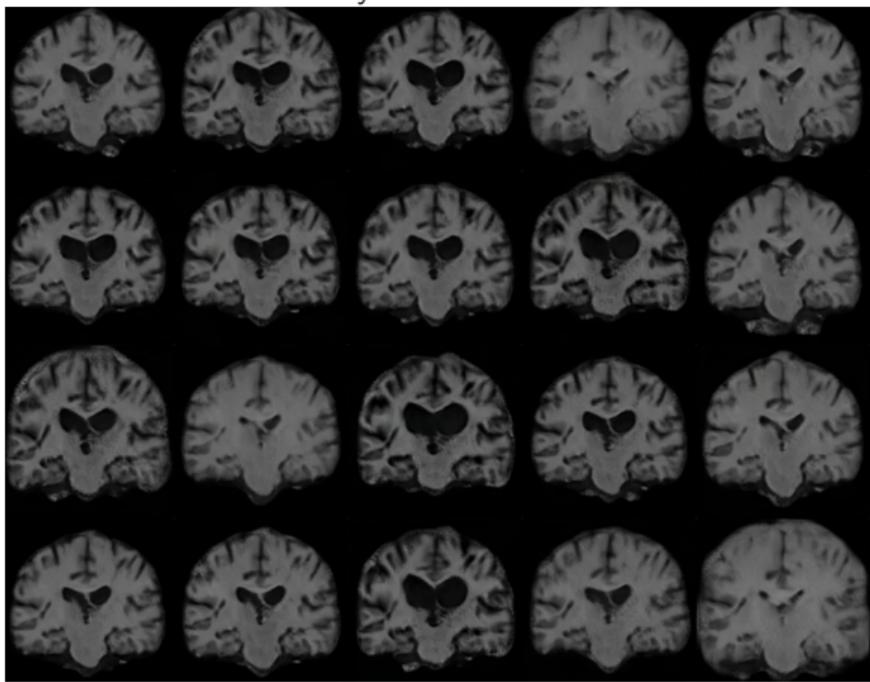
- ▶ StyleGAN2 out-of-the box
- ▶ Trained on MIMIC III, 360k images of 1024x1024 resolution
- ▶ Still some problems to fix:
 - ▶ some ribs look "broken"
 - ▶ bone contours are not always smooth/straight

Preliminary results of StyleGAN2 on brain MRI

Real data



Synthetic data



- ▶ Still out-of-the-box model (StyleGAN2)
- ▶ Trained on 8,000 brain scans (ADNI/OASIS/AIBL/PPMI)
- ▶ Problems to fix:
 - ▶ check if neuroanatomical properties are preserved (e.g. brain/ventricle vols are same)

Preliminary results of StyleGAN2 on microscopy images



- ▶ Trained on 11,000 microscopy slices with pancreatic cancer (MICCAI PANDAS 2020 challenge dataset)

- ▶ GANs have obtained state-of-the-art results on image generation
- ▶ Can generate sharp, realistic images
- ▶ GAN Training is now stable compared to 2-3 years ago, but can take up to 6-7 days (StyleGAN2) on 8 GPUs.
- ▶ Can help solve image reconstruction tasks
- ▶ Many potential applications in medical imaging
- ▶ Recommendations:
 - ▶ Don't build your own, start with a state-of-the-art model (StyleGAN2, BigGAN or Karras, 2020)
 - ▶ Download models already pre-trained to explore their capabilities
 - ▶ When training, initialise weights from another pre-trained model instead of random
- ▶ Keep an eye on other types of generative models (VAEs, auto-regressive, flow) that have other interesting properties (e.g. density estimation)